



Big Data Analytics

Lecture 3:

Computation and Memory Part II

Prof. Dr. Ulrich Matter

11/03/2021

Updates

Schedule update

1. Introduction: Big Data, Data Economy. Walkowiak (2016): Chapter 1.
2. Computation and Memory in Applied Econometrics.
3. Computation and Memory in Applied Econometrics II.
4. Advanced R Programming. Wickham (2019): Chapters 2, 3, 17,23, 24.
5. Import, Cleaning and Transformation of Big Data. Walkowiak (2016): Chapter 3: p. 74-118.
6. Aggregation and Visualization. Walkowiak (2016): Chapter 3: p. 118-127; Wickham et al.(2015); Schwabish (2014).
7. Data Storage, Databases Interaction with R. Walkowiak (2016): Chapter 5.
8. Cloud Computing: Introduction/Overview, Distributed Systems, Walkowiak (2016): Chapter 4.
9. Applied Econometrics with Spark; Machine Learning and GPUs.
10. Project Presentations.
11. Project Presentations; Q&A.

Build groups for group examination

- Teams of 2-3.
- See [Canvas Announcement](#).
- All team members must have a GitHub account.

Group examination: take-home exercises

- Analysis of (big) dataset in R.
- Report in R Markdown.
- Conceptual questions.
- Collaborate, hand-in, feedback via GitHub.

Group projects:

- A simple empirical research question.
- A large (>2GB) data set (of your choice).
 - Get inspired [here](#) and [here](#)
- Implement analysis in R.
- Present results in 6-7 minutes.
 - R-markdown (ioslides/shiny) or R presentation.
- Q&A, Feedback

Group projects:


- Send short **disposition** to ulrich.matter@unisg.ch by end of March.
 - Data set (short description/link)
 - Research question
 - Idea for analysis (statistical approach)
- Have slides ready

Goals for today

1. Understand basics of how to control resource allocation in R.
2. Know the basics of parallel computing in R.
3. Know the basics of efficient memory allocation and virtual memory (in data analytics context).

Recap of Week 2

Components of a computing environment



Components of a computing environment

Why should we care?

Big Data (Analytics)


- Find an efficient (fast) statistical procedure. (Uluru vs OLS example)
- Need to understand how to **make best use of the available resources**, given a specific data analysis task.
 - CPU: Parallel processing (use all cores available)
 - RAM: Efficient memory allocation and usage
 - RAM + Mass Storage: Virtual memory, efficient swapping

Computation and Memory (Part II)

Efficient Use of Resources


1) Parallel processing: CPU/core

- A CPU on any modern computer has several **cores**.
- The OS usually assigns automatically which tasks/processes should run on which core.
- We can explicitly instruct the computer to dedicate N cores to a specific computational task: **parallel processing**.



2) Memory allocation: RAM

- Standard computation procedures happen **in-memory**: data needs to be loaded into RAM.
- Default lower-level procedures to **allocate memory** might not be optimal for large data sets.
- We can explicitly use **faster** memory allocation procedures for a specific big data task.



3) Beyond RAM: virtual memory

- What if we run out of RAM?
- The OS deals with this by using part of the hard disk as **virtual memory**.
- By explicitly instructing the computer how to use **virtual memory for specific big data tasks**, we can speed things up.

Case study: Parallel processing

We start with importing the data into R.

```
url <- "https://vincentarelbundock.github.io/Rdatasets/csv/carData/MplsStops.csv"  
stopdata <- data.table::fread(url) # skipNul avoids running into encoding issues with this data set
```

Case study: Parallel processing

First, let's remove observations with missing entries (NA) and code our main explanatory variable and the dependent variable.

```
# remove incomplete obs
stopdata <- na.omit(stopdata)
# code dependent var
stopdata$vsearch <- 0
stopdata$vsearch[stopdata$vehicleSearch=="YES"] <- 1
# code explanatory var
stopdata$white <- 0
stopdata$white[stopdata$race=="White"] <- 1
```

Case study: Parallel processing

We specify our baseline model as follows.

```
model <- vsearch ~ white + factor(policePrecinct)
```

Case study: Parallel processing

And estimate the linear probability model via OLS (the `lm` function).

```
fit <- lm(model, stopdata)
summary(fit)

##
## Call:
## lm(formula = model, data = stopdata)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.13937 -0.06329 -0.05473 -0.04227  0.97729
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    0.054733   0.005154  10.619 < 2e-16 ***
## white          -0.019553   0.004465  -4.380 1.19e-05 ***
## factor(policePrecinct)2  0.008556   0.006757   1.266  0.2054
## factor(policePrecinct)3  0.003409   0.006483   0.526  0.5990
## factor(policePrecinct)4  0.084639   0.006232  13.582 < 2e-16 ***
## factor(policePrecinct)5 -0.012465   0.006371  -1.956  0.0504 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.254 on 19078 degrees of freedom
```

Case study: Parallel processing

Compute bootstrap clustered standard errors.

```
# load packages
library(data.table)
# set the 'seed' for random numbers (makes the example reproducible)
set.seed(2)

# set number of bootstrap iterations
B <- 10
# get selection of precincts
precincts <- unique(stopdata$policePrecinct)
# container for coefficients
boot_coefs <- matrix(NA, nrow = B, ncol = 2)
# draw bootstrap samples, estimate model for each sample
for (i in 1:B) {

  # draw sample of precincts (cluster level)
  precincts_i <- sample(precincts, size = 5, replace = TRUE)
  # get observations
  bs_i <- lapply(precincts_i, function(x) stopdata[stopdata$policePrecinct==x,])
  bs_i <- rbindlist(bs_i)

  # estimate model and record coefficients
  boot_coefs[i,] <- coef(lm(model, bs_i))[1:2] # ignore FE-coefficients
}
```

Case study: Parallel processing

Finally, let's compute SE_{boot} .

```
se_boot <- apply(boot_coefs,  
                MARGIN = 2,  
                FUN = sd)  
  
se_boot  
  
## [1] 0.004042725 0.004689611
```

Case study: Parallel processing

Parallel implementation...

```
# install.packages("doSNOW", "parallel")
# load packages for parallel processing
library(doSNOW)
# set the 'seed' for random numbers (makes the example reproducible)
set.seed(2)

# get the number of cores available
ncores <- parallel::detectCores()
# set cores for parallel processing
ctemp <- makeCluster(ncores) #
registerDoSNOW(ctemp)

# set number of bootstrap iterations
B <- 10
# get selection of precincts
precincts <- unique(stopdata$policePrecinct)
# container for coefficients
boot_coefs <- matrix(NA, nrow = B, ncol = 2)

# bootstrapping in parallel
boot_coefs <-
  foreach(i = 1:B, .combine = rbind, .packages="data.table") %dopar% {

    # draw sample of precincts (cluster level)
```


Case study: Parallel processing

As a last step, we compute again SE_{boot} .

```
se_boot <- apply(boot_coefs,  
                 MARGIN = 2,  
                 FUN = sd)
```

```
se_boot
```

```
## (Intercept)      white  
## 0.004375874 0.005805687
```

Case study: Memory allocation

```
#####  
# Big Data Statistics: Flights data import and preparation  
#  
# U. Matter, January 2019  
#####  
  
# SET UP -----  
  
# fix variables  
DATA_PATH <- "../data/flights.csv"  
  
# DATA IMPORT -----  
flights <- read.csv(DATA_PATH)  
  
# DATA PREPARATION -----  
flights <- flights[, -1:-3]
```

Case study: Memory allocation

Inspect the memory usage.

```
# SET UP -----  
  
# fix variables  
DATA_PATH <- "../data/flights.csv"  
# load packages  
library(pryr)  
  
# check how much memory is used by R (overall)  
mem_used()  
  
## 1.08 GB  
  
# check the change in memory due to each step  
  
# DATA IMPORT -----  
mem_change(flights <- read.csv(DATA_PATH))  
  
## 611 kB  
  
# DATA PREPARATION -----  
flights <- flights[,-1:-3]
```

Case study: Memory allocation

‘Collect the garbage’...

```
gc()
```

##		used	(Mb)	gc trigger	(Mb)	max used	(Mb)
##	Ncells	1063431	56.8	1757805	93.9	1757805	93.9
##	Vcells	127595415	973.5	213343868	1627.7	211038278	1610.1

Case study: Memory allocation

Alternative approach (via memory mapping).

```
# load packages
```

```
library(data.table)
```

```
# DATA IMPORT -----
```

```
flights <- fread(DATA_PATH, verbose = TRUE)
```

```
## OpenMP version (_OPENMP)          201511
## omp_get_num_procs()                12
## R_DATATABLE_NUM_PROCS_PERCENT    unset (default 50)
## R_DATATABLE_NUM_THREADS           unset
## R_DATATABLE_THROTTLE               unset (default 1024)
## omp_get_thread_limit()             2147483647
## omp_get_max_threads()              12
## OMP_THREAD_LIMIT                   unset
## OMP_NUM_THREADS                    unset
## RestoreAfterFork                   true
## data.table is using 6 threads with throttle==1024. See ?setDTthreads.
## Input contains no \n. Taking this to be a filename to open
## [01] Check arguments
## Using 6 threads (omp_get_max_threads()=12, nth=6)
## NAstrings = [<<NA>>]
## None of the NAstrings look like numbers.
## show progress = 0
## 0/1 column will be read as integer
```

Case study: Memory allocation

Alternative approach (via memory mapping).

```
# SET UP -----
```

```
# fix variables
```

```
DATA_PATH <- "../data/flights.csv"
```

```
# load packages
```

```
library(pryr)
```

```
library(data.table)
```

```
# housekeeping
```

```
flights <- NULL
```

```
gc()
```

```
##           used  (Mb) gc trigger  (Mb) max used   (Mb)
## Ncells  1052467  56.3   1757805   93.9   1757805   93.9
## Vcells 124429667 949.4  213343868 1627.7 211038278 1610.1
```

```
# check the change in memory due to each step
```

```
# DATA IMPORT -----
```

```
mem_change(flights <- fread(DATA_PATH))
```

```
## 35.8 MB
```

Insight from analyzing methods conceptually

- Methods for big data analytics come with an 'overhead'
 - Additional 'preparatory' steps.
 - Only faster than traditional methods if data set has a certain size!

Insight from analyzing methods conceptually

- Methods for big data analytics come with an **'overhead'**
 - Additional 'preparatory' steps.
 - Only faster than traditional methods if data set has a certain size!
- Examples:
 - Parallel processing: Distribute data/task, combine afterwards.
 - `fread`: Memory maps data before actually reading it into RAM.

Beyond memory

- RAM is not sufficient to handle the amount of data to be analyzed...
- **What to do?**

Beyond memory

- RAM is not sufficient to handle the amount of data to be analyzed...
- **What to do?**
- Scale up by using parts of the available Mass Storage (hard-disk) as **virtual memory**

Out-of-memory strategies

- Chunked data files on disk
- Memory-mapped files and shared memory

Out-of-memory strategies

- Chunked data files on disk: `ff`-package
- Memory-mapped files and shared memory: `bigmemory`-package

Chunking data with the ff-package

Preparations

```
# SET UP -----
```

```
# install.packages(c("ff", "ffbase"))
```

```
# load packages
```

```
library(ff)
```

```
library(ffbase)
```

```
library(pryr)
```

```
# create directory for ff chunks, and assign directory to ff
```

```
system("mkdir fdfd")
```

```
options(fftempdir = "ffdf")
```

Chunking data with the ff-package

Import data, inspect change in RAM.

```
##           used  (Mb) gc trigger   (Mb) max used   (Mb)
## Ncells  1086709 58.1   1757805  93.9   1757805   93.9
## Vcells 12891936 983.6  213343868 1627.7 211038278 1610.1
```

```
mem_change(
  flights <-
    read.table.ffdf(file=" ../data/flights.csv",
                    sep=";",
                    VERBOSE=TRUE,
                    header=TRUE,
                    next.rows=100000,
                    colClasses=NA)
)
```

```
## read.table.ffdf 1..100000 (100000)  csv-read=0.34sec ffd-f-write=0.043sec
## read.table.ffdf 100001..200000 (100000)  csv-read=0.361sec ffd-f-write=0.04sec
## read.table.ffdf 200001..300000 (100000)  csv-read=0.366sec ffd-f-write=0.03sec
## read.table.ffdf 300001..336776 (36776)  csv-read=0.14sec ffd-f-write=0.016sec
## csv-read=1.207sec ffd-f-write=0.129sec TOTAL=1.336sec
```

```
## -30 MB
```

Chunking data with the ff-package

Inspect file chunks on disk and data structure in R environment.

show the files in the directory keeping the chunks

```
list.files("ffdf")
```

```
## [1] "clone1664b7fbd953f.ff" "clone1664b9b8cca9.ff" "clone1e7014c0a1cd8.ff"
## [4] "clone1e7015a4f712e.ff" "clone2aea22211d9e1.ff" "clone2aea2360c6703.ff"
## [7] "clone2aea2566ab42d.ff" "clone2aea25e1c1f75.ff" "clone2d49618dbfbf6.ff"
## [10] "clone2d4965ee3349a.ff" "clone2d49664b07745.ff" "clone2d49672b82b88.ff"
## [13] "clone308112a4ca401.ff" "clone308113d044b7c.ff" "clone308113d22fb5f.ff"
## [16] "clone3081149714ed4.ff" "clone399cd5627eb1f.ff" "clone399cd72c6506d.ff"
## [19] "clone399cd78f6c4e6.ff" "clone399cd8b1f075.ff" "clone3c3ef1e38eca1.ff"
## [22] "clone3c3ef4ac46441.ff" "clone3c3ef514956e9.ff" "clone3c3efcb5fb24.ff"
## [25] "ff1664b222c38f0.ff" "ff1664b4d23ee78.ff" "ff1664b4d7f1e3e.ff"
## [28] "ff1e7011754e092.ff" "ff1e7011a76d5a6.ff" "ff1e7017084631e.ff"
## [31] "ff2aea22c3703b9.ff" "ff2aea2664ee33.ff" "ff2aea26b164ce7.ff"
## [34] "ff2d49627cd458.ff" "ff2d49631ca5a34.ff" "ff2d4964237cc21.ff"
## [37] "ff30811207897c4.ff" "ff308115699c1a.ff" "ff30811b3430e.ff"
## [40] "ff399cd1a2ffc0e.ff" "ff399cd1e963877.ff" "ff399cd5477c29.ff"
## [43] "ff3c3ef17300293.ff" "ff3c3ef2229a09b.ff" "ff3c3ef765d6fb7.ff"
## [46] "ffdf1664b11f63957.ff" "ffdf1664b12d379a7.ff" "ffdf1664b16a5a516.ff"
## [49] "ffdf1664b16cc8da8.ff" "ffdf1664b16d72904.ff" "ffdf1664b178a08cd.ff"
## [52] "ffdf1664b17c18654.ff" "ffdf1664b23c24fe0.ff" "ffdf1664b24c84c7c.ff"
## [55] "ffdf1664b25a0763c.ff" "ffdf1664b2663b0d0.ff" "ffdf1664b283091fe.ff"
## [58] "ffdf1664b2b44b5a4.ff" "ffdf1664b2c262d7d.ff" "ffdf1664b2cd1a62.ff"
## [61] "ffdf1664b2e3a4c4e.ff" "ffdf1664b2ed055dd.ff" "ffdf1664b3054a5d7.ff"
```

Memory mapping with **bigmemory**

Preparations

SET UP -----

load packages

`library(bigmemory)`

`library(biganalytics)`

Memory mapping with bigmemory

Inspect the imported data.

```
summary(flights)
```

##	min	max	mean	NAs
## year	2013.000000	2013.000000	2013.000000	0.000000
## month	1.000000	12.000000	6.548510	0.000000
## day	1.000000	31.000000	15.710787	0.000000
## dep_time	1.000000	2400.000000	1349.109947	8255.000000
## sched_dep_time	106.000000	2359.000000	1344.254840	0.000000
## dep_delay	-43.000000	1301.000000	12.639070	8255.000000
## arr_time	1.000000	2400.000000	1502.054999	8713.000000
## sched_arr_time	1.000000	2359.000000	1536.380220	0.000000
## arr_delay	-86.000000	1272.000000	6.895377	9430.000000
## carrier	9.000000	9.000000	9.000000	318316.000000
## flight	1.000000	8500.000000	1971.923620	0.000000
## tailnum				336776.000000
## origin				336776.000000
## dest				336776.000000
## air_time	20.000000	695.000000	150.686460	9430.000000
## distance	17.000000	4983.000000	1039.912604	0.000000
## hour	1.000000	23.000000	13.180247	0.000000
## minute	0.000000	59.000000	26.230100	0.000000
## time_hour	2013.000000	2014.000000	2013.000261	0.000000

Memory mapping with **bigmemory**

Inspect the object loaded into the R environment.

```
flights
```

```
## An object of class "big.matrix"  
## Slot "address":  
## <pointer: 0x5608c09d84d0>
```

Memory mapping with `bigmemory`

- `backingfile`: The cache for the imported file (holds the raw data on disk).
- `descriptorfile`: Metadata describing the imported data set (also on disk).

Memory mapping with bigmemory

Understanding the role of `backingfile` and `descriptorfile`.

First, import a large data set without a backing-file:

```
# import data and check time needed
```

```
system.time(  
  flights1 <- read.big.matrix("../data/flights.csv",  
                              header = TRUE,  
                              sep = ",",  
                              type = "integer")  
)
```

```
##      user  system elapsed  
##  1.022    0.013    1.034
```

```
# import data and check memory used
```

```
mem_change(  
  flights1 <- read.big.matrix("../data/flights.csv",  
                              header = TRUE,  
                              sep = ",",  
                              type = "integer")  
)
```

```
## 736 B
```


Memory mapping with `bigmemory`

Understanding the role of `backingfile` and `descriptorfile`.

Third, re-import the same data set with a backing-file.

```
# remove the loaded file
```

```
rm(flights2)
```

```
# 'load' it via the backing-file
```

```
system.time(flights2 <- attach.big.matrix("flights2.desc"))
```

```
##      user  system elapsed
```

```
## 0.000   0.000   0.001
```

```
flights2
```

```
## An object of class "big.matrix"
```

```
## Slot "address":
```

```
## <pointer: 0x5608bf6f3930>
```

References